

Rapid Scheduling: A Constraint Satisfaction Approach to River Usage

Team: 16677

February 13, 2012

Contents

1 Preliminaries	2
1.1 Introduction	2
1.2 Definitions and Terminology	2
1.3 Assumptions and Simplifications	3
2 The Solution	4
2.1 Constraint Satisfaction	5
2.1.1 Backtracking Search	5
2.2 Usage of High-Level Parameters	7
3 Predictions and Analysis	8
3.1 General Expectations	8
3.2 Results	8
3.3 Analysis	10
4 Final Thoughts	10
4.1 Advantages to the Approach	10
4.2 Detractors from the Approach	11
4.3 Future Extensions	12
4.4 Final Recommendations	13

1 Preliminaries

1.1 Introduction

Floating down a river is easy enough, but imagine scheduling an entire season of rafting trips for hundreds of people. A *scheduling problem* occurs when there are start and end times and start and end locations for a group of people or objects. The schedules of the group must be arranged such that all participants start and stop as desired. Between the start and stop, they must follow any rules stated in the problem.

We consider a scheduling problem for the Big Long River with visitors who desire to raft down it. The goal is to schedule as many visitors as possible in a six month season. Visitors depart from First Launch and finish at Final Exit, 225 miles downstream. They are seeking a relaxing and authentic experience and will become dissatisfied if they do not receive one. While scheduling dissatisfied visitors is allowed, we seek to schedule them so as many visitors are as satisfied as possible.

1.2 Definitions and Terminology

- Big O: (of an algorithm) Also “complexity.” A function relating the time it takes to solve a problem and the inputs of the problem. For instance, a function whose Big O is $O(x)$ computes in linear time.
- Carrying Capacity: The amount of visitors a river can see in a single year, while maintaining the likelihood that all visitors will be scheduled and remain satisfied.
- Constraint Satisfaction Problem (CSP): A type of problem solved by assigning values to variables. These values must be within specified domains. The problem is solved when all variables have values within the specified domains – e.g. they have values that satisfy all constraints.
- Group: See “visitor.”
- Itinerary: The ordered list of campsites a visitor will stay at each night.
- Non-deterministic Polynomial-Time Hard Problem (NP-hard): A class of problems whose solutions, if they exist, behave unpredictably, and are provable in polynomial time – see “Big O.” Constraint Satisfaction Problems are a type of NP-hard problem.
- Visitor: Also “group.” A person or a group of people seeking to travel the Big Long River by raft. Their raft could be powered by oar, or it could be motorized.

Visitors are very irritable and will become unhappy if they see other visitors, don't get their desired water time each day, or don't get scheduled.

- **Satisfaction:** (of a visitor) Visitors seeking to raft the Big Long River start with a satisfaction of 100. This number will drop each time the visitor sees another visitor, or doesn't get their desired water time (too much or too little of it) each day. It is possible for a visitor's satisfaction to drop to zero by these two factors alone. It will be zero if the visitor can't be scheduled.
- **Season:** A six month season is assumed to have 180 days; that is, six months times an average of 30 days per month.
- **Success Rate:** The likelihood a season will have a valid schedule given certain parameters.
- **Tolerance:** The amount of campsites a visitor's travel plan can be altered by. For instance, a tolerance of 1 would allow a visitor desiring to move four campsites each day to move between three and five campsites each day. This is done in order to accommodate for other visitors who may be in the desired spot.
- **Water Time:** Visitors specify how much time they want to spend on the water each day. This value is the same for each day of the trip. A visitor's satisfaction rating will drop if they are forced to spend more or less than this amount of time on the water.
- **X:** The current number of trips taken by visitors down the Big Long River each year.
- **Y:** The number of campsites on the Big Long River. Unless otherwise stated, this was assumed to be 128, as provided by literature[2].

1.3 Assumptions and Simplifications

- **Season Length:** A six month season is assumed to have 180 days; that is, six months times an average of 30 days per month. Seasons are consistent, i.e. there exist no peak visitation months or months of higher river flow.
- **Trips ranging from "6 to 18 nights of camping on the river, start to finish"** are assumed to allow between 7 and 19 days of travel on the river, start to finish.

- Visitors or groups of visitors travel at a constant speed throughout the trip. We define this speed by the type of boat they use. Oar-powered rubber rafts travel at 4 mph, and motorized boats travel at 8 mph.
- No visitor separates into more than one group or combines with other groups during the trip.
- "Contact with other groups of boats on the river" is only counted if one group passes another group or is passed by another group.
- Groups are capable of passing one another at any point (i.e. passing is not restricted by the width of the river or anything else at any time).
- The satisfaction of a group is affected exclusively by the amount of contact with other visitors they have, how much water time they have compared to the desired amount, and whether or not they are actually scheduled.
- Visitors specify only their starting date, boat type, and water time, and only valid values for these parameters. Therefore: They may not start before the start of the season or so close to the end of the season it would be impossible for them to finish; they may not have a water time such that their trip takes less than 7 or longer than 19 days; they may only choose an oar-powered rubber raft or a motorized boat.
- Visitors follow the schedule given to them exactly: no backtracking, skipping, or slowing down.

2 The Solution

Our solution is structured to solve constraint satisfaction problems, specifically of a scheduling nature. Our working definition for such a problem follows: given variables (X), values (D), and constraints (C), find an assignment of value(s) in D for each X such that all constraints in C are satisfied. Generally speaking, a CSP may be defined as a set (X, D, C) . A good example of a CSP is any map coloring problem[1].

We view the problem of scheduling rafting trips as a progressive and dynamic CSP. Each day is a constraint satisfaction problem dependent upon the solution to the previous day's problem. By limiting the values any variable may take on and applying efficient algorithms, the CSP becomes tractable. This allows us to successfully schedule river trips without the need of unreasonable computing power.

2.1 Constraint Satisfaction

Consider that there are exponentially many combinations of visitors to campsites during a season. Many of these combinations, however, are invalid since no two visitors may stay at the same campsite on the same day. This is perhaps the greatest limiting factor in the problem as it applies to all campers. This is an n -ary constraint: a constraint relating n variables. Such constraints are often computationally challenging, but our limited problem space makes the computations relatively painless. We simply pair all campers and detect collisions at any campsite ($O(n^2)$ with n as number of visitors). There are few other constraints involved in solving the problem.

Each visitor is considered to be a variable in the CSP and campsites are considered the values. The domain of each visitor's values varies day to day as the visitor travels through the campsites. Travel distance is derived from a visitor's average speed of 4 or 8 mph, v , multiplied with a uniformly distributed random variable for water time, t . This variable must allow the visitor to complete the trip within 7-19 days. The following equations explicitly describe these relationships:

$$v * t = \text{OptimalDayDistance}$$

$$v * t * 7 \geq L$$

We ensure generated visitors are reasonably parameterized to finish a trip on time. Each visitor is assigned a random uniformly distributed day for departure, $dDay$, such that

$$v * t * 7 + dDay \leq S$$

A visitor's maximum travel speed will allow them to finish before the close of the season.

Visitors are generated en masse using uniformly distributed variables. It is possible to load a file containing visitor information so the program will work on problems of natural rather than simulated origin. Once visitors are generated and basic parameters are set (discussed in **Section 2.2**) the program begins attempting a solution through our solver.

2.1.1 Backtracking Search

In order to solve a CSP, one must navigate many combinations for variables and values. Backtracking search algorithms are often used to assign values to variables[4]. We use such an algorithm to assign campsites to our visitors on a daily basis. On any given day

d , the solver looks to see which visitors are set to depart that day and adds them to a list, *toDepart*. Then, all visitors left on the river from the previous day are added to *toDepart*. This creates a list of all visitors requiring movement for day d . In order to determine appropriate movements, visitors are assigned campsites one-by-one. Travel between campsites is analagous to distance traveled for the day. Following a visitor-campsite pairing, the day's assignments are checked to see if this pairing is consistent. If the pairing doesn't conflict with existing pairings, it is permanently added to the day's assignments. We then proceed to the next visitor.

When pairings fail, the algorithm attempts new pairings until success is achieved or no more pairings exist. In this case, failure occurs; the algorithm backtracks to the previous variable; undoes that pairing; and attempts a new pairing for this prior variable. The entire process is structured as a depth-first-search. When a day's assignment is complete, the algorithm creates a new CSP for the next day and executes again. **Figure 1** highlights the major steps involved in this process with a simple example. The algorithm begins by assigning site 1 to visitor A , denoted by the diamond below and to the left of the tree-top. Assigning site 1 to B is tried but found to fail. Instead the algorithm assigns site 2 to B (rightmost circle). At this point, the day would “end” and the algorithm would begin scheduling the next day. In the event the next day fails, the algorithm returns to a previous level, as seen in the example. In this case, the previous level is the previous day.




Figure 1: An example of the backtracking algorithm's attempt to schedule visitors A and B with campsites 1 and 2.

Our solver seeks maximal satisfaction of the travelers. That is, the possible campsites a visitor may travel to are presented to the solver in order of ideal travel distance. While not every visitor is guaranteed to travel their desired distance, the algorithm attempts to satisfy it.

We also calculate the number of encounters between visitors. After the itineraries are created, each visitor's campsite listing is checked to determine how many visitors passed and were passed by this specific visitor. By combining the amount of variation from ideal travel distance and encounters with other visitors, we obtain the total effect on satisfaction. The solver relies heavily on several parameters to relax the problem, making solutions more attainable.

2.2 Usage of High-Level Parameters

Many values tend to be static. These include river length, number of days in the season, and number of campsites. While each of these values is capable of drastically affecting the outcome of a given attempt, this effect is limited by a well-defined range (see **Section 1.3**). There are likely optimal values, but finding these would likely place our problem into the realm of intractable problems. In an attempt to capture the complexity of this system, we have introduced a unique variable, *tolerance*.

Tolerance is designed to create a range for the distance a visitor may travel in a given day. This variable is straightforward at the outset but provides several neat nuances for additional complexity in the solution. The introduction of a value range allows for a larger domain of solutions and variability. We desire such flexibility to account for the real-world uncertainty of directing human beings. Tolerance extends as an aggregate for many of the instances where we might expect variability —such as inclement weather, cancellations, or slow visitors— without having to introduce additional parameters. Tolerance is a satisfying opportunistic variable for allowing more visitors to be scheduled. It allows for the rafting season to be enjoyed by more individuals.

It is important to note a lack of tolerance often creates problems with no solutions. With visitors locked between one of two travel distances, scheduling becomes virtually impossible for all but the smallest number of visitors.

3 Predictions and Analysis

3.1 General Expectations

Before creating the model we developed a few “common sense” expectations of the results.

- Increasing the visitor size would make it harder to schedule each visitor at the desired time. The more visitors one is trying to schedule, the faster the schedule fills up.
- As the number of visitors on the river during a season increases, it becomes increasingly likely visitors will pass each other. Passing is a necessity if the visitors are traveling at different daily speeds. Consider a visitor with a daily speed of 8 mph leaving a few days after a visitor traveling at 4 mph.
- As tolerance increases, so will the ease of scheduling each visitor. The greater the tolerance, the greater the likelihood of finding a site to stay at if other visitors are also in the area.
- An increase in tolerance will cause a decrease in overall happiness. An increase in tolerance will allow visitors to travel less than or more than their desired time each day. Thus decreasing their overall satisfaction.
- Assigning more visitors to motor boats (those that travel at 8mph) will increase the success rate. Visitors that travel faster will spend fewer days on the river, leaving more space for other visitors.
- Seasons with less variability in visitor speeds (e.g. seasons where all visitors travel at 4mph, or where all visitors travel at 8mph) will result in greater satisfaction. Visitors traveling at the same speed are less likely to pass each other or to want to stay at the same campsite on the same night. Greater variation in group speeds will increase the frequency of these conflicts.

3.2 Results

The results of the model tended to be in line with our expectations (**Section 3.1**). Our baseline values for parameters are in **Section 1.3**. For seasons with a 2-site tolerance, the success rate was seen to decrease as the number of visitors increased (**Figure 2 Spark 1**).

The success rate was 100% for seasons with 5, 50, and 100 visitors. Once the number of visitors per season increased past 100, the success rate began decreasing at an increasing rate. The success rate dropped from 100% to 91% between 100 and 300 campers, then from 91% to 25% between 300 and 500 campers. For seasons with more than 500 campers, the success rate leveled off, appearing to have an asymptote just above zero percent. The success rate for a season with 600 visitors was five percent.

Similarly, the average happiness of each of the visitors decreased as the number of visitors increased (**Figure 2 Spark 2**). Once the number of visitors per season increased past 100, the happiness rate began decreasing at an increasing rate. The happiness dropped from 97.9% to 89.4% between 100 and 300 campers, then from 89.4 percent to 24.5% between 300 and 500 campers. For seasons with more than 500 campers, the happiness rate leveled off, appearing to have an asymptote just above zero percent. The happiness for a season with 600 visitors was 4.9%.

For seasons with a 1-site tolerance, the effect of visitor size was seen to increase (**Figure 2 Spark 3**). Both the success rate and happiness of visitors began dropping with a visitor size of only 50 people. The success rate dropped from 100% to 96% when visitor sized increased from 5 to 50 visitors per season. By just 250 visitors, both the success rate and the happiness were zero percent.

The tolerance for which campsite visitors stayed at each night had a positive effect on success rate of the scheduler and happiness of the visitors (**Figure 2 Spark 4**). For seasons with 200 visitors, a tolerance of zero campsites had a success rate of 0% and happiness of zero percent. Increasing the tolerance to one campsite resulted in a success rate of 52% and a happiness of 4.9%. Increasing the tolerance to two campsites had the greatest effect, raising success rate to 97.8% and happiness to 95.9%. With any greater tolerances, the success rate was one-hundred percent, and the happiness leveled off at 97.8%.

For seasons with 500 visitors, a greater tolerance was needed to achieve similar results (**Figure 2 Spark 5**). A tolerance of 0 campsites had a success rate of 0; a tolerance of 1 campsite had a success rate of only 2.6%; 2 campsites, 28%; 3 campsites, 90%.



Figure 2: 1: Success Rate vs. Visitors 2: Satisfaction vs. Visitors 3: Satisfaction vs. Visitors 4: Tolerance vs. Success Rate and Satisfaction 5: Tolerance vs. Success Rate and Satisfaction

The success rate was only 100% with a tolerance of greater than 4 campsites, and the happiness was 98%.

Assigning more visitors to motorboats (those that travel at 8 mph) was seen to increase the success rate. For a homogeneous season of 8 mph visitors, the success rate was 98.2% and the happiness 74%. At the number of 4 mph boats increased, both success rate and happiness steadily decreased. With half rowboats and half motorboats, the success rate was 92% and the happiness 72%. With a 9:1 ratio of rowboats to motorboats, the success rate of the season was only 89.9%. However, there is a jump in success rate to 99% once all visitors are traveling at 4 mph. This jump in successes is not accompanied by a jump in happiness; with all rowboats, the happiness is only 70%.

3.3 Analysis

As defined in our definitions section, carrying capacity is the amount of visitors a river can see in a single year while maintaining the likelihood that all visitors will be scheduled and remain satisfied. With a tolerance of 2 campsites (**Figure 3**), we found the river could manage 100 visitors per season with a 100% success rate and a satisfaction of at least 97%. It is worth noting that doubling the capacity to 200 visitors per seasons yields a success rate of 99.5% and a satisfaction still greater than 97%. Once the number of visitors per season increases past 200, both success rate and tolerance begin to drop rapidly. With a tolerance of 1 campsite, we found the river could only manage about five visitors per season with a success rating of 100%. The average satisfaction of these visitors is 97.7%. After the number of visitors per season increases above 50, both satisfaction and success rate begin dropping rapidly as seen in **Figure 4**.

Is our value of 300 visitors, which was reasonably successful and maximal, realistic? Assuming that it is possible to have 60 individuals in a visitor group, we find that our model allows for 18,000 people in a season to travel down the Big Long River. Our sources suggest this is realistic for comparable rivers (i.e. the Colorado River)[3].

4 Final Thoughts

4.1 Advantages to the Approach

Throughout, we have seen that a constraint satisfaction approach to scheduling visitors to a river is a natural fit. This approach allows for a logical approach to the scheduling process, and for fairly quick solutions. In most cases, solutions were found on mid-grade hardware in less than ten minutes. We also found that the structure of our solution




Figure 3: Groups per Season vs. Success Rate and Satisfaction with Tolerance 2

in both code and at a higher level lent itself well to the adjustment of any number of parameters. Moreover, there exists a great deal of flexibility once one begins to define “fuzzy values” that we can assign to the variables. Given more time, extending the model with these fuzzy values would likely lead to better assignments and even more interesting results.

4.2 Detractors from the Approach

river variation perfection of visitors vary start date In our assumptions (**Section 1.3**) we note that the visitors are basically perfect. This, of course, is not true. Visitors will probably vaguely follow the schedule given them: due either to inclement weather, rudeness, or physical constraints. It is not immediately apparent how to deal with this shortcoming other than simply accepting that not all variables can be accounted for. Padding the schedules with “free days” on which any amount of travel lag may be recuperated.

The use of a consistent river in our model is less detrimental to the overall performance and accuracy. A realistic river would include areas where passing other groups is impossible, visitors travel faster or slower, and the potential for divisions of the river. These could all be built in as constraints on the campsites and their assignment to




Figure 4: Groups per Season vs. Success Rate and Satisfaction with Tolerance 1

visitors. Should we have more time, this would likely be incorporated into the model.

4.3 Future Extensions

Given more time there are several additions that would improve the technical and practical performance of the model:

- The addition of a filtering technique to detect early failure in the backtracking search. This is often performed as what is called “arc-consistency”. Regretfully, it is difficult to implement for our n-ary constraints.
- Utilizing a fuzzy constraint for departure days would allow groups to have a window in which they might leave. This is likely to increase the number of possible visitors, considering that it is nearly a direct analog to tolerance.
- The schedule might be modified to include days of rest for campsites on a rotating basis. This would not aid in adding more visitors (in fact it would decrease the success rate), but it would assist in land management practices and encourage a healthy use of the river ecosystem.

4.4 Final Recommendations

Designating an exact number of visitors to raft down a river each year is a complex problem that may not have only one solution. Managers will need to be assertive in designating a tolerance, allotting departure days and boat types to visitors, and instructing visitors to follow their schedules carefully. Given certain parameters, however, we recommend a tolerance of 2 campsites, and about 300 visitors per season. Depending upon the current number of visitors each season, X , **this would mean increasing the number of visitors by $300 - X$** . Note that those 300 visitors may be *groups*, not necessarily individuals. A tolerance of 2 campsites permits a large number of visitors per season while maintaining acceptable visitor satisfaction. Any greater tolerance would create a sharp decrease in satisfaction as visitors are forced to travel unfortunate distances; any smaller tolerance greatly decreases the chance of a successful schedule for any large number of visitors. Any more than 300 visitors is likely to drop both the chance of a successful season and the happiness of visitors.

References

- [1] R. Dechter and A. Dechter, *Belief maintenance in dynamic constraint networks in proc. of aaai-88*, 1988.
- [2] Will Hansen, *Whitewater Campsites colorado–grand canyon*, 2009.
- [3] Catherine A. Roberts Randy Gimblett, *Integrating gis and agent based modeling techniques for understanding social and ecological processes*, Oxford Press, 2000.
- [4] Stuart J. Russell and Peter Norvig, *Artificial intelligence: A modern approach*, Prentice Hall, 2010.